
Conversational Music Retrieval with Synthetic Data

Megan Leszczynski^{*†1,2} Ravi Ganti^{*2} Shu Zhang^{*2} Krisztian Balog²

Filip Radlinski² Fernando Pereira Arun Tejasvi Chaganty^{*2}

¹ Stanford University ² Google Research

mleszczy@cs.stanford.edu

{gmravi,shzhang,krisztianb}@google.com

{filiprad,pereira,arunchaganty}@google.com

Abstract

Users looking for recommendations often wish to improve suggestions through broad natural language feedback (e.g., “How about something more upbeat?”). However, building such *conversational* retrieval systems requires conversational data with rich user utterances paired with slates of items that cover a diverse range of preferences. This is challenging to collect scalably using conventional methods like crowd-sourcing. We address this problem with a new technique to synthesize high-quality dialog data by transforming the domain expertise encoded in curated item collections into corresponding item-seeking conversations. The method first generates a sequence of hypothetical slates returned by a system, and then uses a language model to introduce corresponding user utterances. We apply the approach on a dataset of curated music playlists to generate 10k diverse music-seeking conversations. A qualitative human evaluation shows that a majority of these conversations express *believable* sequences of slates and include user utterances that *faithfully* express preferences for them. When used to train a conversational retrieval model, the synthetic data yields up to a 23% relative gain on standard retrieval metrics compared to baselines trained on non-conversational and conversational datasets.

1 Introduction

Recent work has made significant advances in the ability of retrieval systems to understand natural language queries and non-textual content (e.g., images, audio) [11, 12, 22]. However, standard retrieval systems still struggle to retrieve items for ambiguous queries (e.g., “Music for focusing”), where the query’s meaning may depend on the user or their context (e.g., writing, meditating). This motivates *conversational retrievers*, which allow the user to provide natural language feedback (e.g., “How about something more upbeat?”) to steer the system to retrieve the items they are looking for.

Conversational retrievers are challenging to build because they require conversational training data that pairs multiple turns of rich and diverse natural language feedback with retrieved items. One possibility is to generate data in a crowd-sourced Wizard-of-Oz setup [2, 14, 18, 24]: here, one person acts a user looking for items, while another acts as a wizard that recommends items given the user’s queries and feedback. A key limitation of this approach is that both people need some domain expertise: the wizard to find relevant items to suggest, and the user to provide varied and meaningful feedback. For many domains, few people have this expertise, leading to shallow conversations.

^{*}Core contributor

[†]Work done while at interning at Google.

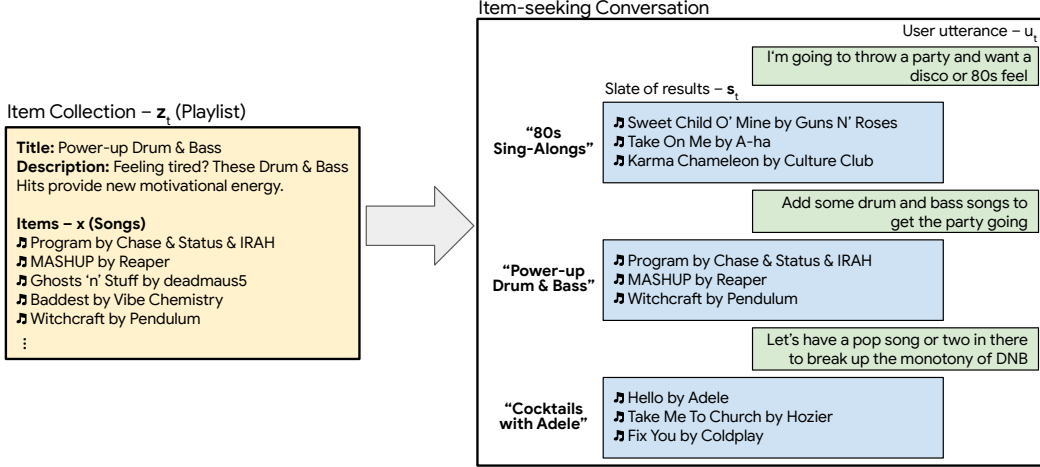


Figure 1: Curated item collections (e.g., music playlists) group items (songs) into coherent collections and provide useful metadata about them (e.g., playlist titles and descriptions). We propose a new method that leverages this domain expertise to create item-seeking conversations that contain user utterances with rich and diverse feedback paired with item slates that reflect this feedback.

In contrast, curated item collections (e.g., playlists, recipe books) are widely available on the internet. These collections capture the domain expertise of their creators who pick a coherent set of items. Moreover, they often include detailed metadata (e.g., titles, descriptions and tags) with attributes a user may refer to in their preferences (e.g., “upbeat music”, “healthy recipes”). Motivated by these properties, we ask: can we leverage curated item collections to generate conversations?

There are two key features missing in item collections: (1) multiple turns with slates of items to recommend to the user, and (2) user utterances describing their feedback for each slate. We solve these problems in two steps: First, we observe that, in an ideal conversation, each turn should make a coherent change that brings the user closer to their target slate. Following Göpfert et al. [8], we represent items, collections, and slates as vectors in a shared embedding space, and generate a sequence of slates by pivoting around item collections towards the target slate. Second, we use a dialog inpainting language model [5] to generate conversational user utterances that express preferences for each slate by prompting the model with metadata from corresponding item collections (Figure 1).

We use this approach to create a dataset of 10,000 synthetic music-seeking conversations that cover a breadth of domain expertise, from conversations about Japanese pop music to those about electro-swing music. In a qualitative evaluation, we find a majority of the data contain believable sequences of slates with user utterances that faithfully express preferences for them. We also evaluate this data quantitatively by measuring its impact on training conversational music retrievers. Our proposed approach yields up to a 23% relative Hits@10 gain compared to baselines trained on non-conversational and conversational datasets.

Related work. Traditional recommender systems [15, 16] use matrix factorization techniques to personalize results from large quantities of user log data. In contrast, conversational or interactive recommender systems [3] can reduce their dependence on logs by allowing users to interact with them and provide direct feedback. However, because of a lack of conversational training data, conversational recommenders are often trained using reinforcement learning techniques [3, 17, 28, 30] or using supervised learning on scripted dialogue flows [9]. A notable exception is the ReDial dataset [18], a large conversational movie recommendation dataset collected by paid experts in a Wizard-of-Oz setting. Our work also benefits from recent advances in conversational systems [4, 6, 26, 27, 27, 29, 29], retrieval modeling [13, 25], and content modeling [10, 11, 19, 21].

2 From Curated Item Collections to Item-Seeking Conversations

Curated item collections contain substantial domain expertise: they not only group items (songs) $x \in \mathcal{X}$ into coherent collections (playlists) $z \subset \mathcal{X}$, but also provide valuable metadata about each

collection $\phi(\mathbf{z})$ (playlist titles and descriptions). However, they lack two key features present in an item-seeking conversation: a sequence of *slates* of items $\mathbf{s}_t \subset \mathcal{X}$ presented to the user in each turn t (instead of a fixed collection of items), and corresponding user utterances u_t that express a preference for each slate (instead of the non-conversational metadata). We address these problems by using a dataset of item collections \mathcal{Z} to first generate sequences of slates and then generate user utterances for them.

Generating sequences of slates. We are guided by following desirable properties on the slates a user might see in an ideal conversation: (P.1) a slate, \mathbf{s}_t , should maintain *continuity* with its previous turn’s \mathbf{s}_{t-1} (if the user first asked for party music, they are more likely to ask for pop music than meditation music); (P.2) the change between subsequent slates should be *coherent*, corresponding to natural language feedback from the user—we expect that an item collection \mathbf{z}_t approximates this change; and finally, (P.3) each turn should bring the user *closer to their target slate* \mathbf{s}^* (if the user ultimately wants good party music, we expect slates to include more high energy music in later turns).

To realize these properties, we follow Göpfert et al. [8] and represent items x , item collections \mathbf{z}_t and slates \mathbf{s}_t as vectors (\tilde{x} , $\tilde{\mathbf{z}}_t$ and $\tilde{\mathbf{s}}_t$) in a shared embedding space \mathbb{R}^d . We use P.1 and P.2 to model $\tilde{\mathbf{s}}_t$ as a linear combination: $\tilde{\mathbf{s}}_t = \alpha\tilde{\mathbf{s}}_{t-1} + \beta\tilde{\mathbf{z}}_t$, and use P.3 to choose α and β to minimize the distance to $\tilde{\mathbf{s}}^*$.³ We sample $\tilde{\mathbf{s}}_0$ and $\tilde{\mathbf{s}}^*$ from \mathcal{Z} to ensure that sequences start and end in coherent slates. Empirically, we found that sampling $\tilde{\mathbf{z}}_t$ from shrinking neighborhoods of $\tilde{\mathbf{s}}^*$ works well⁴; we plan to further explore this in future work. Finally, we can retrieve the items in \mathbf{s}_t using the item neighbors of $\tilde{\mathbf{s}}_t$.

Generating user utterances. Once we have a sequence of slates, we need to generate corresponding user utterances: suppose that \mathbf{s}_{t-1} contained “80s Sing-Alongs” songs and \mathbf{s}_t “Power-Up Drum & Bass” songs; u_t should contain feedback like “Now *drum and bass songs* to get the party going”. To achieve this, we use a dialog inpainter [5]—a T5-based language model trained to predict missing utterances in a conversation. We set up the conversation with system utterances that describe each slate \mathbf{s}_t using the corresponding collection’s metadata $\phi(\mathbf{z}_t)$ —e.g., “I’ve added songs described as *sick party drum and bass songs*.”—and use a dialog inpainter to fill in the “missing” user utterances.⁵

Application to conversational retrieval. Our ultimate goal is to build a conversational item retriever where users can interact with the system by providing feedback over multiple turns. Given a user utterance u_t and the history of previous utterances and slates $\mathbf{H}_t = (u_1, \mathbf{s}_1, \dots, u_{t-1}, \mathbf{s}_{t-1})$, a conversational retriever predicts a new slate of items that ranks the user’s target \mathbf{s}^* highest.⁶ We model the task by learning a *ranking function* $\rho_t : \mathcal{X} \rightarrow \mathbb{R}$ using a dual encoder architecture [7, 13, 20], which has shown to be effective on similar tasks. Dual encoders independently embed queries q and targets x into dense vectors, and compute the ranking function using cosine similarity: $\rho(x; q) = \text{embed}(x) \cdot \text{embed}(q)$. We follow prior work [5, 25] and use item metadata to provide (u_t, \mathbf{H}_t) as a query, and sample items from \mathbf{s}_t to be the target (see Appendix B for model input details).

3 Evaluation

Our evaluation aims to validate two claims: (1) our approach generates high quality synthetic data, and (2) training on the synthetic data improves retrieval performance. We use our approach to generate 10,000 music-seeking conversations from 19,156 expert-curated playlists (\mathcal{Z}) from a proprietary dataset. Each conversation consists of six turns, where each turn t has a user utterance u_t and slate \mathbf{s}_t . The retrieval corpus \mathcal{X} consists of 381,331 songs from the same dataset.

3.1 Dataset Quality

We begin by qualitatively evaluating the synthetic data. One author manually rated over 100 turns on a three point scale according to the following questions: (1) how believable is the turn, with respect to the target?, (2) how well-phrased is the user utterance?, and (3) how well does the user utterance

³ α and β can be solved in closed-form when the vectors in the embedding space are of unit-norm.

⁴See Appendix A.1 for further details on the embedding space and sampling procedure.

⁵See Appendix A.2 for examples of the templates used to generate the system utterances.

⁶We assume that the user’s target is fixed throughout the conversation.




Figure 2: *Left*: Qualitative evaluation of over 100 turns from the synthetic data. *Right*: Retrieval performance on the test set of the synthetic data.

match their preference? We want sequences of slates that are realistic, user utterances that resemble those of real users, and user utterances that faithfully represent the preferences for the slates. We find that the majority of turns are high quality on each of the dimensions, suggesting that the synthetic data is sufficiently high quality to use for training data (Figure 2).

3.2 Retrieval Performance

We now measure the quantitative impact of the synthetic data when used to train conversational retrievers. We first describe the experimental setup. See Appendix C for more details.

Evaluation Dataset. As there is no existing data for the conversational music retrieval task, we evaluate on a test set of the synthetic data, which consists of 856 conversations. We report a standard retrieval metric, Hits@10, by comparing the song ranking at each turn with the target slate, excluding songs that have been seen in the history up to that turn. Hits@10 is 1 iff any of the top-10 retrieved songs are labeled relevant. All models use the same conversation history from the dataset, as opposed to building conversation history using model predictions from previous turns.

Model implementation. We initialize our dual encoder from a pretrained T5 1.1-Base checkpoint [23] and finetune on the synthetic data by randomly sampling a turn t for each conversation, and using $q_t = (u_t, \mathbf{H}_t)$ as a query and a randomly sampled song from s_t as a target. We train the model using a contrastive loss with in-batch negatives. At inference time, we precompute the song embeddings and use nearest neighbor search to retrieve the top 10 songs for each query q_t .

Baselines. We compare against two dual encoders trained on variants of the query: NonConversational, where the query q is the playlist title and seed songs from the playlist, and NoHistory, where query $q \stackrel{\text{def}}{=} u_t$ for a randomly sampled turn t . We also compare against two dual encoders trained on variants of the synthetic dataset: RandomSequence, where the slates s_t are equal to a randomly sampled playlist z_t at each turn t , and UserTemplates, where the user utterances are templated instead of generated by the inpainter. The training target is a randomly sampled song x from the playlist (NonConversational) or from s_t (all others). While the training sets differ between the baselines, all baselines are evaluated on the same set of queries, including the conversation history, at test time.

Figure 2 (right) compares our model to baselines. We observe that our model achieves up to a 23% relative gain over the best baseline on Hits@10 averaged over turns. The gains over the UserTemplates model suggest it is useful to train on conversational user utterances rather than templated user utterances. Moreover, the gains over the NoHistory model suggest that training with a conversation history is also useful. Overall, the results suggest that our model is better able to use natural language feedback to improve retrieval. As future work, we are interested in scaling up the qualitative evaluation of the synthetic data and evaluating the retrieval performance on real conversational item retrieval datasets.

4 Discussion

We introduced a general technique to convert item collections into synthetic conversations and demonstrated the benefits of building conversational retrieval models trained on such data for the

music domain. In the future, we plan to more extensively evaluate our models and explore how to incorporate user information—either from user logs or explicit user intent—to further personalize conversational retrieval systems. Finally, while language models are known to propagate the harmful biases in their training data [1], by synthetically generating data we have the opportunity to filter and mitigate these biases—this is an important direction for future work.

References

- [1] Rishi Bommasani, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, Erik Brynjolfsson, Shyamal Buch, Dallas Card, Rodrigo Castellon, Niladri Chatterji, Annie Chen, Kathleen Creel, Jared Quincy Davis, Dora Demszky, Chris Donahue, Moussa Doumbouya, Esin Durmus, Stefano Ermon, John Etchemendy, Kawin Ethayarajh, Li Fei-Fei, Chelsea Finn, Trevor Gale, Lauren Gillespie, Karan Goel, Noah Goodman, Shelby Grossman, Neel Guha, Tatsunori Hashimoto, Peter Henderson, John Hewitt, Daniel E. Ho, Jenny Hong, Kyle Hsu, Jing Huang, Thomas Icard, Saahil Jain, Dan Jurafsky, Pratyusha Kalluri, Siddharth Karamcheti, Geoff Keeling, Fereshte Khani, Omar Khattab, Pang Wei Koh, Mark Krass, Ranjay Krishna, Rohith Kuditipudi, Ananya Kumar, Faisal Ladhak, Mina Lee, Tony Lee, Jure Leskovec, Isabelle Levent, Xiang Lisa Li, Xuechen Li, Tengyu Ma, Ali Malik, Christopher D. Manning, Suvir Mirchandani, Eric Mitchell, Zanele Munyikwa, Suraj Nair, Avaniika Narayan, Deepak Narayanan, Ben Newman, Allen Nie, Juan Carlos Niebles, Hamed Nilforoshan, Julian Nyarko, Giray Ogut, Laurel Orr, Isabel Papadimitriou, Joon Sung Park, Chris Piech, Eva Portelance, Christopher Potts, Aditi Raghunathan, Rob Reich, Hongyu Ren, Frieda Rong, Yusuf Roohani, Camilo Ruiz, Jack Ryan, Christopher Ré, Dorsa Sadigh, Shiori Sagawa, Keshav Santhanam, Andy Shih, Krishnan Srinivasan, Alex Tamkin, Rohan Taori, Armin W. Thomas, Florian Tramèr, Rose E. Wang, William Wang, Bohan Wu, Jiajun Wu, Yuhuai Wu, Sang Michael Xie, Michihiro Yasunaga, Jiaxuan You, Matei Zaharia, Michael Zhang, Tianyi Zhang, Xikun Zhang, Yuhui Zhang, Lucia Zheng, Kaitlyn Zhou, and Percy Liang. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- [2] Paweł Budzianowski, Tsung-Hsien Wen, Bo-Hsiang Tseng, Iñigo Casanueva, Stefan Ultes, Osman Ramadan, and Milica Gašić. MultiWOZ - a large-scale multi-domain Wizard-of-Oz dataset for task-oriented dialogue modelling. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018.
- [3] Konstantina Christakopoulou, Filip Radlinski, and Katja Hofmann. Towards conversational recommender systems. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*, pages 815–824, 2016.
- [4] Konstantina Christakopoulou, Alex Beutel, Rui Li, Sagar Jain, and Ed H Chi. Q&R: A two-stage approach toward interactive recommendation. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 139–148, 2018.
- [5] Zhuyun Dai, Arun Tejasvi Chaganty, Vincent Y Zhao, Aida Amini, Qazi Mamunur Rashid, Mike Green, and Kelvin Guu. Dialog inpainting: Turning documents into dialogs. In *Proceedings of the 39th International Conference on Machine Learning*, 2022.
- [6] Jianfeng Gao, Michel Galley, and Lihong Li. Neural approaches to conversational ai. *Foundations and Trends® in Information Retrieval*, 2019.
- [7] Daniel Gillick, Sayali Kulkarni, Larry Lansing, Alessandro Presta, Jason Baldridge, Eugene Ie, and Diego Garcia-Olano. Learning dense representations for entity retrieval. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 528–537, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/K19-1049. URL <https://aclanthology.org/K19-1049>.
- [8] Christina Göpfert, Yinlam Chow, Chih-Wei Hsu, Ivan Vendrov, Tyler Lu, Deepak Ramachandran, and Craig Boutilier. Discovering personalized semantics for soft attributes in recommender systems using concept activation vectors. In *Proceedings of the ACM Web Conference 2022, WWW ’22*, page 2411–2421, 2022.

- [9] Javeria Habib, Shuo Zhang, and Krisztian Balog. Iai moviebot: A conversational movie recommender system. In *Proceedings of the 29th ACM International Conference on Information & Knowledge Management*, 2020.
- [10] Curtis Hawthorne, Andrew Jaegle, Cătălina Cangea, Sebastian Borgeaud, Charlie Nash, Mateusz Malinowski, Sander Dieleman, Oriol Vinyals, Matthew Botvinick, Ian Simon, Hannah Sheahan, Neil Zeghidour, Jean-Baptiste Alayrac, Joao Carreira, and Jesse Engel. General-purpose, long-context autoregressive modeling with perceiver AR. In *Proceedings of the 39th International Conference on Machine Learning*, 2022.
- [11] Qingqing Huang, Aren Jansen, Joonseok Lee, Ravi Ganti, Judith Yue Li, and Daniel P. W. Ellis. Mulan: A joint embedding of music audio and natural language. In *Proceedings of the the 23rd International Society for Music Information Retrieval Conference (ISMIR)*, 2022.
- [12] Chao Jia, Yinfei Yang, Ye Xia, Yi-Ting Chen, Zarana Parekh, Hieu Pham, Quoc Le, Yun-Hsuan Sung, Zhen Li, and Tom Duerig. Scaling up visual and vision-language representation learning with noisy text supervision. In *Proceedings of the 38th International Conference on Machine Learning*, 2021.
- [13] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6769–6781, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.550. URL <https://aclanthology.org/2020.emnlp-main.550>.
- [14] J. F. Kelley. An iterative design methodology for user-friendly natural language office information applications. *ACM Trans. Inf. Syst.*, 1984.
- [15] Yehuda Koren, Robert Bell, and Chris Volinsky. Matrix factorization techniques for recommender systems. *Computer*, 42(8):30–37, 2009.
- [16] Yehuda Koren, Steffen Rendle, and Robert Bell. Advances in collaborative filtering. *Recommender systems handbook*, pages 91–142, 2022.
- [17] Wenqiang Lei, Xiangnan He, Yisong Miao, Qingyun Wu, Richang Hong, Min-Yen Kan, and Tat-Seng Chua. Estimation-action-reflection: Towards deep interaction between conversational and recommender systems. In *Proceedings of the 13th International Conference on Web Search and Data Mining*, pages 304–312, 2020.
- [18] Raymond Li, Samira Ebrahimi Kahou, Hannes Schulz, Vincent Michalski, Laurent Charlin, and Chris Pal. Towards deep conversational recommendations. *Advances in neural information processing systems*, 31, 2018.
- [19] Ilaria Manco, Emmanouil Benetos, Elio Quinton, and György Fazekas. Learning music audio representations via weak language supervision. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 456–460. IEEE, 2022.
- [20] Jianmo Ni, Chen Qu, Jing Lu, Zhuyun Dai, Gustavo Hernández Ábrego, Ji Ma, Vincent Y. Zhao, Yi Luan, Keith B. Hall, Ming-Wei Chang, and Yinfei Yang. Large dual encoders are generalizable retrievers. *arXiv preprint arXiv:2112.07899*, 2021.
- [21] Aaron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew Senior, and Koray Kavukcuoglu. Wavenet: A generative model for raw audio. *arXiv preprint arXiv:1609.03499*, 2016.
- [22] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 8748–8763. PMLR, 18–24 Jul 2021. URL <https://proceedings.mlr.press/v139/radford21a.html>.

- [23] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020. URL <http://jmlr.org/papers/v21/20-074.html>.
- [24] Tsung-Hsien Wen, David Vandyke, Nikola Mrksic, Milica Gasic, Lina M Rojas-Barahona, Pei-Hao Su, Stefan Ultes, and Steve Young. A network-based end-to-end trainable task-oriented dialogue system. *arXiv preprint arXiv:1604.04562*, 2016.
- [25] Shi Yu, Zhenghao Liu, Chenyan Xiong, Tao Feng, and Zhiyuan Liu. Few-shot conversational dense retrieval. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR '21*, page 829–838, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450380379. doi: 10.1145/3404835.3462856. URL <https://doi.org/10.1145/3404835.3462856>.
- [26] Yichi Zhang, Zhijian Ou, and Zhou Yu. Task-oriented dialog systems that consider multiple appropriate responses under the same context. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 9604–9611, 2020.
- [27] Yongfeng Zhang, Xu Chen, Qingyao Ai, Liu Yang, and W Bruce Croft. Towards conversational search and recommendation: System ask, user respond. In *Proceedings of the 27th acm international conference on information and knowledge management*, pages 177–186, 2018.
- [28] Xiaoxue Zhao, Weinan Zhang, and Jun Wang. Interactive collaborative filtering. In *Proceedings of the 22nd ACM international conference on Information & Knowledge Management*, pages 1411–1420, 2013.
- [29] Jie Zou, Yifan Chen, and Evangelos Kanoulas. Towards question-based recommender systems. In *Proceedings of the 43rd international ACM SIGIR conference on research and development in information retrieval*, pages 881–890, 2020.
- [30] Lixin Zou, Long Xia, Yulong Gu, Xiangyu Zhao, Weidong Liu, Jimmy Xiangji Huang, and Dawei Yin. Neural interactive collaborative filtering. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 749–758, 2020.

Appendix

A Data Generation Details

We provide additional implementation details for generating sequences of slates and generating user utterances described in Section 2.

A.1 Slate Sequence Generation

We generate six turns for each conversation using hyperparameters that we found generated reasonable sequences in preliminary experiments. We randomly sample the starting point \tilde{s}_0 from the neighborhood of the target \tilde{s}^* using a neighborhood of size $k = 256$, using the softmax of similarity to the target as the sampling distribution. To ensure that the starting point is not too close to the target, we also require that \tilde{s}_0 is not in the top 50 neighbors of \tilde{s}^* . We then sample \tilde{z}_t from the neighborhood of \tilde{s}^* , decreasing k each turn until $k = 1$ on the last turn to ensure that the conversation reaches the user’s target slate. Specifically, we sample from neighborhood sizes of $\{128, 64, 32, 16, 1\}$ and use the softmax over the similarity to the previous turn’s slate \tilde{s}_{t-1} as the sampling distribution to encourage each \tilde{z}_t to be relevant to the previous slate. Finally, we sample the top 20 items that are neighbors of \tilde{s}_t to retrieve the items in the slate s_t for each turn. We are interested in understanding the sensitivity of our approach to these hyperparameters in future work.

Preference	System Utterance Template
<i>More</i>	<p>Of course! Let me also add some songs described as <i>\$description</i>. How can I improve the vibe?</p> <p>Got it, I will also add some songs described as <i>\$description</i>. Describe how I can improve the vibe.</p> <p>Sure, I will also add some songs described as <i>\$description</i>. How can I improve this playlist now?</p> <p>Definitely, let me also add more songs described as <i>\$description</i>. What else can I do to make this playlist better?</p> <p>Ok, let me go ahead and also add some songs described as <i>\$description</i>. How can I improve this playlist specifically?</p>
<i>Less</i>	<p>Got it, let me remove some songs described as <i>\$description</i>. How would you like to improve this playlist now?</p> <p>Sure, I will remove some songs described as <i>\$description</i>. How can I improve the vibe?</p> <p>Ok, let me remove some songs described as <i>\$description</i>. What else can I do to make this playlist better?</p> <p>Of course! I will remove some songs described as <i>\$description</i>. How can I improve this playlist now?</p> <p>Sounds good. Let me remove songs described as <i>\$description</i>. Describe how I can improve the vibe.</p>

Table 1: Examples of templates used to instantiate system utterances. We fill in *\$description* with the description for playlist \mathbf{z}_t .

A.2 Utterance Generation

After generating sequences of slates, we use a dialog inpainter to generate corresponding user utterances. The dialog inpainter takes as input a sequence of system utterances. We create the system utterances by filling in templates with metadata for each item collection \mathbf{z}_t . We include examples of the templates used for the system utterances to generate music-seeking conversations in Table 1. We support two types of preferences—*more* or *less*—which we determine based on the value of the weight β in the linear combination (i.e., positive β means *more*).

B Model Inputs

We describe how we represent the inputs to the non-conversational and conversational dual encoders.

B.1 Non-conversational dual encoder

We use a non-conversational dual encoder to learn the joint embedding space to generate the sequences of slates in Section 2. We also use this model as one of the baselines in Section 3.2. Here, we discuss how we represent the inputs, the query q and the targets x , to the dual encoder.

The query q is a textual representation of an item collection \mathbf{z} , generated using item collection metadata and item metadata for a sample of "seed" items in \mathbf{z} . The target x is a textual representation of a random sampled item in \mathbf{z} , generated using the item metadata.

Concretely, for music retrieval, the query input is as follows:

```
playlist_title [SEP] seed_song_title by artist_names from album_title
```

In our experiments, we use five seed songs and concatenate a description of each seed song to the end of the query with a separator token. Similarly, the target input is defined as:

```
song_title by artist_names from album_title
```


B.2 Conversational dual encoder

We now discuss how we represent the inputs to the conversational dual encoder. Recall that we define the conversation history at turn t as $\mathbf{H}_t = (u_1, \mathbf{s}_1, \dots, u_{t-1}, \mathbf{s}_{t-1})$. For the conversational dual encoder, query $q \stackrel{\text{def}}{=} (u_t, H_t)$, where u_t is the user utterance at turn t . The target x is a sampled item from the slate \mathbf{s}_t .

To construct the query, we observe that the user utterances u are already text. We then represent slate \mathbf{s} textually using the metadata associated with the top- k items from \mathbf{s} . Finally, we concatenate the textual representation of each term in \mathbf{H}_t with the utterance u_t . Note that we use reverse concatenation (i.e., the last turn occurs at the beginning of the query and the first turn occurs at the end of the query), so that the position of the last utterance u_t is independent of the turn t . The targets use the same representation as the non-conversational model (Section B.1).

Specifically, for the conversational music retriever, the query input, at time t is as follows:

utterance _{t} [SEP] song_description _{$t-1$} [SEP] utterance _{$t-1$} ... [SEP] song_description₁ [SEP] utterance₁

Where "utterance" represents the user utterance u and "song_description" represents a textual description of the song sampled from the slate \mathbf{s} that matches the target input encoding. In our experiments we use the top 3 songs from the slate \mathbf{s} at each turn and concatenate the song descriptions with a separator token (shown with the top 1 song for space).

C Retrieval Experiment Details

We provide additional details for the retrieval results in Section 3.2.

C.1 Data preprocessing

One failure mode of the dialog inpainter is occasionally copying text from the input prompts. We want to reduce these cases as it can lead to user utterances that are less conversational and less realistic. To account for this, we filter the test split of the synthetic data to remove conversations that have a high substring overlap between a user utterance and the system utterances (i.e., longest common substring is greater or equal to 75 characters). This filters about 14% of conversations, resulting in a test set of 856 conversations.

C.2 Baselines

We include the templates used to instantiate the user utterances for the `UserTemplates` baseline in Table 2. We assign preferences as described in Appendix A.2.

Preference	User Utterance Template
<i>More</i>	Add some songs described as <i>\$description</i> . Can you also add some songs described as <i>\$description</i> . I want more songs that can be described as <i>\$description</i> . Can I have some songs described as <i>\$description</i> . More <i>\$description</i> .
<i>Less</i>	Fewer songs described as <i>\$description</i> . Please no songs described as <i>\$description</i> . I don't want songs described as <i>\$description</i> . Remove songs described as <i>\$description</i> . Less <i>\$description</i> .

Table 2: Examples of templates used to instantiate *user* utterances for the `UserTemplates` baseline. We fill in *\$description* with the description for playlist \mathbf{z} .

C.3 Training details

We finetune all models (including baselines) from a T5.1.1-Base checkpoint for 500k steps with constant learning rate 0.001, dropout rate 0.1, and batch size 512. We use TPUv3 chips to train the non-conversational model and TPUv4 chips to train our model and all other baselines.

C.4 Evaluation details

We use a query token length of 1024 tokens and a target token length of 128 tokens and preprocess the input as described in Section B.2 for all models. We use the validation set of the synthetic data to select the best checkpoint for our model and baselines. We run inference using the full retrieval corpus, and select the checkpoint that achieves the highest Hits@10 averaged over turns. We report Hits@10 results using the best checkpoint for each model over the test set of the synthetic data.